Vemana-Base: Telugu → English evaluation

Actual offline inference | 3 short pilot notes + the same 100 long synthetic Telugu notes

Model and license

udgamlabs/Vemana-Base; 278,214,144 parameters; Apache 2.0. The model card identifies it as a Telugu base model that continues text rather than following instructions. Its trained context is 1,024 tokens. No fine-tuning or model substitution was performed.

Observed results

Measurement	Result
Long notes attempted	100 (302–317 words each)
Passed basic English/completion checks	0/100 — not accuracy
Rejected outputs	100
Context-budget failures	0
Copied garbage-example phrase	100/100
Source-topic keyword present	10/100 — weak diagnostic only
100-note run / peak process RAM	3.05 min / 2.19 GiB
Execution	CPU FP32; long-note batches of 8

Examples from the short pilots — model output, not corrected

Water outage → direct prompt:

A woman who was involved in a murder case in Kurnool district, was arrested by the police. The woman was arrested by the police … [excerpt]

Water outage → one-example prompt:

Garbage has not been collected in the street for one week. The resident requests the Garbage Department. Story first published: Monday, May 24, 2009, … [excerpt]

Pension delay → one-example prompt:

Garbage has not been collected in the street for one week. The resident requests the government to give the pension to the poor.

Assessment and limits

The direct water prompt produced an unrelated crime story; the few-shot prompt copied its garbage example and invented details. English-looking output is not evidence of source-faithful summarization. The model loads and generates Telugu, but this trial does not support using it for the requested summarizer. Task-specific supervised adaptation would need a separate evaluation.

All source text was retained without truncation; plain completions were used because no chat template is supplied. Outputs and hashes are in results.json/all_outputs.csv. The synthetic dataset and prompt differences prevent a controlled accuracy comparison with Qwen. No GPU or AWS instance was used for this trial.

Model card/license: huggingface.co/udgamlabs/Vemana-Base